

CrackID: A Benchmark for Crack Re-Identification, and the Limits of Appearance Matching

Anonymous Submission
Paper ID — *under review*

Abstract—Automated maintenance systems turn photographs into tracked defect records, so a defect re-photographed across visits must be re-identified: the new photo must match the specific physical crack’s record, or open a new one. Despite mature crack *detection*, we know of no dataset or benchmark for crack *re-identification*, and the natural approach — crop the region and match embeddings — imports an assumption from person/vehicle re-ID that may not transfer: that a crack has stable, locally distinctive appearance. We introduce CRACKID: 140 photographs over 17 walls with 212 crack identities, released with an explicit annotation-quality report and with every metric’s chance level stated. Each wall was captured in one continuous walk-around, so the benchmark measures *within-visit viewpoint robustness*; a synthetic multi-visit protocol (warp plus grown or corrupted appearance) is reported as a secondary check, and re-identification across real sessions is a hypothesis here, not a result. Benchmarking nine matchers across four families, we find that on the decision a system actually makes — *same crack or not*, at a threshold — none exceeds chance by more than 0.054 pairwise F_1 , two achieve nothing, and at a validation-calibrated threshold eight of the nine span only -0.012 to $+0.012$ around chance while LoFTR holds a margin of 0.046. These are point estimates on an effective sample of 19 identities over 11 walls; the wall-level intervals are wide (Sec. VIII), so the claim is the structure — a band at the floor with one exception — not the individual digits. Two controls localise the difficulty: giving an embedder the surrounding wall raises its open-set rate $0.168 \rightarrow 0.404$, but *erasing the crack* from that view costs almost none of it (0.393) — such a method retrieves on location, not crack appearance. A geometric reference baseline — register on wall texture with cracks masked out, Chamfer-match projected crack masks under global assignment — reaches $\text{DIR@FAR}=0.1$ of 0.768 against 0.429 for the best crop method and 0.000 for a control that exploits its coverage without examining a crack, so the open-set gap is attributable to matching, not partial coverage; like the F_1 figures these are point estimates on the same 19-identity sample, and the claim is the 0.35 gap as a structure, not a precise spread (Sec. VIII). This is not yet deployment evidence: the geometric system answers only 48% of pairs, fails categorically on two smooth-plaster walls, and its apparent advantage is entangled with homography-based annotation provenance. On the synthetic revisit check, which is the paper’s only multi-visit evidence, OSNet+context instead wins ranking and open-set performance; the geometric advantage therefore should not be read as a demonstrated cross-visit solution.

Index Terms—crack re-identification, dataset, benchmark, building maintenance, structural inspection, homography, geometric alignment, annotation quality

I. INTRODUCTION

Automated building-maintenance systems convert photographs into tracked defect records: a re-photographed defect must be *re-identified* — matched to its record, or given a new

one. Wall cracks are a common and difficult case, because the same crack is re-photographed from different distances, angles and lighting, and (Sec. III) grows.

Crack analysis is a mature *detection* and *segmentation* task, and its datasets — CrackSeg9k [1], OmniCrack30k [2] — carry no cross-observation identity labels, because the task they serve needs none. We know of no benchmark for crack *re-identification*; the natural approach, crop-and-match as in person and vehicle re-ID [3], [4], assumes a crack has stable, locally distinctive appearance, and whether it does has never been measured. We introduce CRACKID — 140 photographs, 17 walls, 212 crack identities — to make that measurement (Fig. 1). Capture is one continuous walk-around per wall, so what is measured is *within-visit* viewpoint and scale correspondence — an upper bound on multi-visit performance; we add a synthetic revisit protocol as a venue-controlled secondary check (Sec. VII-D).

Appearance buys almost nothing on the decision. On ranking the nine matchers separate as expected (Rank-1 0.352–0.654 against 0.336 chance); on pairwise F_1 — the *same-crack-or-not* decision a record makes at a threshold — all nine land between 0.000 and 0.054 above the 0.305 floor, two (ORB, DeiT) at exactly zero, and at a validation-calibrated threshold eight of the nine span only -0.012 to $+0.012$ around chance while LoFTR holds a real margin of 0.046.

Two controls say where the difficulty is not. Giving an embedder the surrounding wall raises OSNet’s open-set rate $0.168 \rightarrow 0.404$; *erasing the crack* from that view costs almost none of it (0.393) — such a method retrieves on the wall, and the cue that works, even implicitly, is *location*.

A geometric baseline makes location explicit. Registering on non-crack texture, projecting crack masks into a common frame, and Chamfer-matching under global assignment reaches $\text{DIR@FAR}=0.1$ of 0.768 against 0.429 (best crop) and 0.404 (best context), answering 48% of pairs; a COVERAGE-ONLY control scores 0.000 on the same metric, so the open-set gap is matching, not alignment coverage. This is a conditional result, not a deployment-ready system: two test walls have zero registration success, and the annotation and baseline both use homography-based alignment.

The revisit check reverses the ranking story. The real data are one continuous walk-around per wall, minutes apart, not separate visits. Under the synthetic warp-plus-appearance protocol, OSNet+context beats the hybrid at R@1 (0.826/0.671 versus 0.275/0.274) and open-set rate (0.739/0.663 ver-




Fig. 1. **The appearance ceiling.** Pairwise F_1 for nine crop-scope matchers spanning four method families on CRACKID test. Every method falls inside the band 0.305–0.359 — width 0.054 — pressed against the 0.305 chance floor ($2p/(1+p)$ for this pool), irrespective of family, era, or compute budget. The geometric reference baseline (red) leaves the band. Ranking metrics separate the same methods over a 0.30 range; the decision metric does not.

sus 0.343/0.342) for growth/appearance change; the hybrid’s small assignment- F_1 advantage is not evidence that geometry solves real multi-visit re-identification.

Contributions: (i) CRACKID with an explicit annotation-quality report — 45% of identities were formed by an automatic merge, reported rather than assumed (Sec. III-B); (ii) a nine-method benchmark against stated chance levels, with crop+context and crack-erased controls (Sec. V); (iii) a synthetic revisit protocol isolating appearance change from damage evolution — the only multi-visit evidence this paper has (Sec. VII-D); (iv) a geometric reference baseline and a protocol for this task’s failure modes — coverage-aware metrics, unanswerable queries, an admission gate, wall-resampled confidence (Sec. VI, IV, VII).

II. RELATED WORK

Crack analysis is overwhelmingly detection and segmentation [5], [6]; its datasets annotate cracks within single images and the large consolidations merging [1], [2] inherit that scope, and multi-view pipelines [7] aggregate within one inspection without instance correspondence. Structural-health monitoring does compare cracks across time with machinery close to ours [8]–[11], but identity is always *given by construction* (fixed camera, marked location, operator at a viewpoint): the subject is one known crack, the question width or extent, and no dataset supports method comparison on the association decision itself.

Object re-identification [4] with open-set identification [12] is the formal setting we adopt; keypoint matching [13], [14], learned correspondence (SuperGlue [15], LoFTR [16]), and region embeddings [3], [17]–[20] apply directly to crack crops and our benchmark measures all of them. Vehicle-damage assessment already re-identifies damage with OSNet in curated settings [21]; our results suggest the crop scope is itself the limiting choice. The move our baseline makes — align two captures, then reason about what is present in one and absent in

the other — is change detection [22], [23]; planar registration is classical [24]–[26] with photometric alignment (ECC [27]) as the fallback. Our contribution is not a new estimator but the observation that the alignment target and the matching target should be *disjoint*: the homography is driven by the wall precisely because the crack is what we seek to identify.

III. CRACKID

A. Capture and composition

We photographed 17 exterior walls at one site over four half-days with two consumer handsets, yielding 140 photographs at 4080×3072 . The walls span two surface families that behave very differently under keypoint matching: stippled render (densely textured) and smooth painted plaster, where the crack is frequently the only structure in frame. Walls split disjointly into 6 validation walls (66 photographs, 66 identities) and 11 test walls (74 photographs, 146 identities). Statistical power is governed by identities visible in two or more photographs — only these form a query with an answer: 18 in validation, 19 in test. Each wall was captured in one continuous walk-around (consecutive frames a median of 3 s apart), so the dataset measures **viewpoint and scale correspondence** within one visit — an upper bound on multi-visit performance, not yet a measurement of it.

The adjacent frame is always an answer. For 100% of the 280 answerable test queries the nearest correct gallery entry is the adjacent frame. Ranking numbers are therefore near-duplicate retrieval; Sec. VII sweeps the exclusion gap to show the structure is not what produces the result.

B. Annotation, and why it needs auditing

One click point per crack is propagated across consecutive photographs through an ORB homography and one-to-one assignment; identities whose components appear to be fragments of one physical crack are then merged. The merge is where risk concentrates: a wrongly joined merge converts a hard negative into a free positive and inflates *every* method’s mAP together, so it cannot be caught by comparing methods. We therefore report its extent rather than assume its correctness: 95/212 **identities (45%) were formed by a merge** (48% on test), bridging a median maximum gap of 228 px; the adjusted Rand index against the pre-merge labelling is 0.080 — the merge restructured the partition rather than refining it. We release the instrumentation for an audited error rate (per-identity provenance, a stratified sample with contact sheets, Wilson scoring) rather than the audit itself; the verdicts are not yet filled in, so we claim no sampled error rate.

IV. BENCHMARK PROTOCOL

A query is one annotated crack instance in one photograph; the gallery is the query’s own wall, so hard negatives share material and lighting; the query’s own photograph and week are excluded, leaving $n_Q = 280$ scored queries on test. Three choices deliberately favour no method, including ours:

- **Unanswerable queries are retained.** A query whose identity has no gallery counterpart can contribute false

alarms and never successes — exactly the open-set case a maintenance system faces.

- **Thresholds are calibrated on validation.** Every method reports an oracle-thresholded F_1^* (generous to all) and a deployable $F_1^{\text{@v}}$ thresholded on held-out walls.
- **Missing scores are charged, not skipped.** A pair a method cannot score counts as predicted-negative and ranks last; coverage is not independent of the label (Sec. VII-B).

A. Metrics

We report Rank-1, Rank-5 and mAP, pairwise F_1 , and assignment F_1 (a one-to-one matching per image pair, the deployed operation). The headline metric is the open-set **detection-and-identification rate at 10% false alarm** (DIR@FAR=0.1): a known query is correct only if its top-ranked match is above threshold and correct; an unknown query is a false alarm whenever anything clears threshold.

Every metric is reported against its chance level, because two are far from zero: pairwise F_1 has a structural floor $2p/(1+p) = 0.305$ on this pool, and Rank-5 chance is 0.767; Rank-1 and DIR@FAR chance are 0.336 and 0.038 (closed form, plus 200 random matrices through the same evaluation code). Confidence resamples **walls**, not queries — queries from one wall succeed or fail together — and with 11 test walls the interval is wide; we report its width as the honest statement.

V. THE APPEARANCE CEILING

A. Setup

Nine crop-scope matchers — handcrafted keypoints (SIFT [13], ORB [14]); learned correspondence (SuperGlue [15], LoFTR [16]); general embeddings (ViT-B/16 [17], DeiT-S [18], CLIP ViT-B/32 [19], YOLOv8); a re-identification architecture (OSNet [3]) — run under an identical protocol, inputs and scoring.

B. Result: appearance buys almost nothing on the decision

On the decision a record actually makes — *same crack or not*, at a threshold — no crop method clears chance by more than 0.054 (excess over the 0.305 floor at the oracle threshold, Table I: 0.000–0.054, two methods exactly at the floor, eight within 0.052), across two decades of methods and two orders of compute (≈ 1 ms to 8.6 s). At a deployable (validation-calibrated) threshold eight of the nine span only -0.012 to $+0.012$ around chance; **LoFTR is the exception**, $F_1^{\text{@v}} = 0.351$ ($+0.046$), and the most expensive crop method. Ranking does not transfer to deciding: the nine separate over a 0.302 range on Rank-1, yet OSNet ties for best Rank-1 (0.654) while sitting 0.004 above chance, and SuperGlue (fifth on Rank-1) has the best crop open-set rate (0.429). Ordering a gallery well and keeping a record are not the same capability. The collapse after validation calibration is consistent with threshold instability under a small, wall-clustered sample: methods can order candidates differently while their score

distributions do not support a stable same-crack/not-same-crack boundary on held-out walls. This is an interpretation of the observed pattern, not a claim that calibration itself is the sole cause.

C. Is it the crack, or the wall behind it?

Two controls separate identity cue from field of view: a *crop+context* condition (wider crop) and a *crack-erased* condition (the crack inpainted out of that widened view). Context helps a lot — OSNet’s Rank-1 rises $0.654 \rightarrow 0.775$, mAP to 0.595, open-set rate $0.168 \rightarrow 0.404$ (CLIP’s $0.143 \rightarrow 0.311$; F_1^* $0.309 \rightarrow 0.354$) — **but erasing the crack costs almost none of it**: OSNet retains Rank-1 0.711, mAP 0.574 and open-set 0.393; CLIP retains Rank-1 0.600, mAP 0.502, open-set 0.286. Both erased views still outrank every tight-crop method: the retrieved cue is the wall, i.e. location. Pairwise F_1 is unmoved by context (0.346–0.361). We read the aggregate as a property of the data — methods with no shared inductive bias land near the floor, and wherever a crop method improves, the improvement is not the crack. Nine methods are a sample, not proof; the result establishes that the obvious moves buy 0.000–0.054 above chance on the decision — reason to change the formulation, not the encoder.

VI. A GEOMETRIC REFERENCE BASELINE

The ceiling motivates a change of scope. A crack is a mark at a fixed location on a rigid, near-planar surface that carries texture the crack does not; we identify a crack by *where it is on its wall* once the two photographs are aligned. We release this as a deliberately simple *reference baseline* for CRACKID, with no learned component beyond the shared crack segmenter.

A. Alignment driven by non-crack texture

A wall is planar, so the homography $x' \sim Hx$ mapping photograph I_a into I_b ’s frame is the correct model. The design decision that matters is **what drives the fit**: we dilate the crack masks by 15px and exclude those pixels from detection, so H comes from wall texture alone. This is not an optimisation but what makes the projected position an independent identity cue — were cracks to participate, a long crack could dominate the correspondence set and pull the fit toward the wrong crack. Detection runs on copies downsampled to 1600px under CLAHE equalisation; we use SIFT (up to 8000 features), a 0.75 ratio test, and MAGSAC at 4px requiring 25 inliers.

B. A cascade for two surface families

One front end does not serve both surfaces, and the failure is categorical: on the smooth-plaster walls, crack-masked SIFT succeeds on *exactly zero* pairs. We try three in order: (i) **SIFT on masked wall texture**, for stippled render; (ii) **SIFT with cracks restored**, triggered only when masking starves detection below 300 keypoints (on plaster the crack is the only feature in frame), accepting the circularity stage (i) avoids, and recording which front end produced each result; (iii) **ECC**, direct photometric alignment over a $320 \rightarrow 640 \rightarrow 1280$ px pyramid, on the large-scale shading that survives on flat paint.

TABLE I

MAIN RESULT, TEST SPLIT, $n_Q = 280$ QUERY INSTANCES OVER 11 WALLS. ALL METHODS SEE IDENTICAL INPUTS AND ARE SCORED BY IDENTICAL CODE. F_1^* IS ORACLE-THRESHOLDED ON TEST; ΔF_1^* IS ITS EXCESS OVER THE 0.305 CHANCE FLOOR ($2p/(1+p)$, SEC. IV); $F_1^{\text{@}v}$ IS THRESHOLDED ON THE HELD-OUT VALIDATION WALLS AND IS THE DEPLOYABLE NUMBER. *scored* IS THE FRACTION OF VALID QUERY–GALLERY PAIRS RECEIVING A FINITE SCORE; UNSCORED PAIRS ARE CHARGED AS PREDICTED-NEGATIVE AND THE QUERY IS STILL COUNTED (SEC. VII-B). COST IS WALL-CLOCK SECONDS PER QUERY. THE COVERAGE-ONLY REFERENCE ROW ANSWERS YES TO EVERY REGISTERED PAIR WITHOUT EXAMINING A CRACK AND IS THE LEVEL EVERY OTHER ROW MUST BE READ AGAINST; CHANCE LEVELS FOR EVERY COLUMN ARE GIVEN IN SEC. IV. BEST CROP-SCOPE VALUE PER COLUMN IN **BOLD**.

Method	Family	R@1	mAP	DIR@FAR=0.1	F_1^*	ΔF_1^*	$F_1^{\text{@}v}$	scored	s/query
<i>Reference level</i>									
COVERAGE-ONLY	geometry	0.515	0.533	0.000	0.544	0.239	0.544	0.48	—
<i>Crop scope</i>									
ORB [14]	keypoint	0.352	0.364	0.018	0.305	0.000	0.305	1.00	0.038
SIFT [13]	keypoint	0.432	0.417	0.036	0.313	0.008	0.305	1.00	0.086
SuperGlue [15]	learned	0.562	0.455	0.429	0.353	0.048	0.305	1.00	5.826
LoFTR [16]	learned	0.549	0.440	0.336	0.359	0.054	0.351	1.00	8.639
YOLOv8 backbone	embedding	0.561	0.453	0.157	0.357	0.052	0.308	1.00	0.217
ViT-B/16 [17]	embedding	0.654	0.511	0.286	0.331	0.026	0.317	1.00	0.193
DeiT-S [18]	embedding	0.621	0.497	0.250	0.305	0.000	0.293	1.00	0.058
CLIP ViT-B/32 [19]	embedding	0.604	0.471	0.143	0.307	0.002	0.306	1.00	0.050
OSNet [3]	re-ID	0.654	0.520	0.168	0.309	0.004	0.309	1.00	0.053
<i>Full-image scope</i>									
REG+CHAMFER (ours)	geometry	0.952	0.886	0.768	0.656	0.351	0.582	0.48	27.92

Every candidate H is sanity-checked (reject mirrored/folded quadrilaterals, area scale outside $[1/16, 16]$, collapsed sides); its inlier RMS sets downstream tolerances.

C. Projection and Chamfer agreement

Crack instances are connected components; pixels from I_a are projected by H and compared with instances of I_b . Chamfer distance, not IoU, is the criterion: cracks are a few pixels wide, so a 3px registration error drives IoU to zero for a correctly aligned pair. The reported score is the **symmetric mean Chamfer distance** in pixels

$$D(a, b) = \frac{1}{2} \left[\frac{1}{|a|} \sum_{p \in a} d(p, b) + \frac{1}{|b|} \sum_{q \in b} d(q, a) \right], \quad (1)$$

over the union bounding box, with cheap rejection when boxes are far apart. We also report a scale-free alternative — **Chamfer coverage**, the symmetric fraction of pixels within τ — because Eq. (1) is a per-pair pixel distance whose scale differs per pair, so a single global threshold makes ranking partly measure calibration (Sec. VII-C).

D. Explicit non-answers, and the licensed negative

Correspondence is a one-to-one assignment per image pair with an explicit rejection threshold, not a per-query argmax — two cracks must not both claim the same crack; unmatched instances are reported as *disappeared* or *newly appeared*. The Hungarian solve is part of *evaluation*, applied identically to every method’s score matrix. Specific to this formulation is the evidence a **licensed negative** can carry: with I_a projected into I_b ’s frame, the system observes a region visible and empty of the recorded crack — positive evidence of absence,

categorically stronger than the low similarity a crop embedding produces. A projected crack falling outside the gallery photograph is a confident rejection, counted separately from unscorable pairs; this is the mechanism behind the open-set gap (Sec. VII-B).

VII. RESULTS

A. Main comparison and the open-set gap

Relocating identity into the wall raises every metric: Rank-1 $0.654 \rightarrow 0.952$, mAP $0.520 \rightarrow 0.886$, and as excess over the COVERAGE-ONLY reference the ranking gains are $+0.437/+0.353$, while pairwise $F_1^* = 0.656$ is $+0.112$ over 0.544, not the $+0.351$ a comparison against the crop floor would suggest. The **open-set result is the one number coverage does not explain**: DIR@FAR=0.1 reaches 0.768 against 0.429 for the best crop method (SuperGlue), 0.404 for the best context control, 0.038 for chance, and exactly **0.000** for COVERAGE-ONLY. Crop methods cannot decline — they return a similarity to the nearest neighbour, which on self-similar hairline cracks is often high — while an aligned pair supports the *licensed negative*: the query projects into the gallery frame, where its region is visible and empty. Compute does not explain the gap: SuperGlue’s 0.429 costs 5.8 s/query, OSNet’s 0.168 costs 0.053 s.

B. Partial coverage, gap structure, and failures

Every crop method scores 100% of valid pairs, the geometric method 48%; unscorable pairs are charged as predicted negatives. For **ranking** that is a penalty: censoring a complete matrix drops Rank-1 $0.422 \rightarrow 0.27$ at 34% coverage. For **pairwise** F_1 it is a subsidy: the scorable subset retains 100% of positives and 36.8% of negatives, so prevalence rises $0.180 \rightarrow 0.373$ and the floor with it, $0.305 \rightarrow 0.544$ —

COVERAGE-ONLY scores 0.544 pairwise, mAP 0.533, Rank-1 0.515, above OSNet, by answering every registered pair without examining a crack. Much of the ranking and pairwise- F_1 margin is therefore *which pairs align*; the open-set metric is the exception and the result we lead with.

Excluding the adjacent frame, then the two and three nearest, leaves Rank-1 flat (OSNet 0.654 $\rightarrow$ 0.648) with excess over chance stable in 0.24–0.32: the near-duplicate structure is real but not what produces the numbers. It is also important for interpretation: the adjacent frame answers all 280 queries, whereas the gap-3 sweep leaves only 105 answerable queries, so the headline ranking measures short-baseline correspondence rather than broad temporal re-identification. Registration success is strongly wall-dependent (97/301 pairs, 32.2%: walls 07/08/12 100%, wall 09 52%, wall 11 83%, walls 13–16 between 3% and 29%, walls 10/17 at 0%). A logistic fit over all 301 pairs makes the mechanism plain: sharpness dominates (median 133.8 registered vs 7.2 failed; wall-clustered Wald $z = 5.1$, $p = 3 \times 10^{-7}$), frame gap far less ($z = -2.8$, $p = 0.005$). The independent-pairs z s (8.2, -5.0) overstate the contrast because failure is categorical by wall, so we resample walls, not queries — and the leave-one-wall-out spread in Rank-1 is large, so no fine ranking among baselines is claimed. One admission gate (variance-of-Laplacian ≥ 10) applied identically rejects 21 of 74 test captures and nearly doubles registration, 32.2% $\rightarrow$ 56.1%.

C. Ablations and cost

Chamfer distance versus coverage. The scale-free coverage score motivated in Sec. VI is not a free improvement: it drops mAP 0.886 $\rightarrow$ 0.646 and Rank-1 0.952 $\rightarrow$ 0.849, the largest movement any change produces; DIR and $F_1^{\text{@v}}$ barely move — coverage gives a threshold decision, not a ranking. **Masking.** Disabling crack exclusion leaves ranking and pairwise agreement practically unchanged (to three decimals) but raises DIR@FAR 0.768 $\rightarrow$ 0.825: masking is not what buys the reported accuracy. This is expected when the test set contains no deliberate case where crack pixels create a spurious registration: masking buys evidence independence, not observed accuracy on this collection, and should be treated as a safeguard awaiting a harder stress test. **ECC.** Disabling ECC (repaired after a pipeline bug silently left it inert) gives Rank-1 0.920, mAP 0.862, coverage 0.50 against 0.952, 0.886, 0.48; assignment F_1 rises 0.769 $\rightarrow$ 0.784 when enabled; open-set is unchanged.

Cost. The abstract’s more-than-two-orders-of-magnitude comparison at 27.9 s/query is a query-level charging convention, not a fundamental pair-level compute gap. At 27.9 s/query the baseline charges more than two orders of magnitude above OSNet, but that per-query figure is a charging convention, not an intrinsic one: registration is per image *pair* — one homography serves every crack on a wall — and amortised over the 301 test image pairs the same pass runs ≈ 26 s/pair (7817 s total), after which any later crack query on the pair pays only the Chamfer pass, not the homography. That pair-

Protocol	Method	n_q	R@1	mAP	DIR@FAR=0.1	$F_1^{\text{@v}}$	F_1^a	scored
A (growth)	Hybrid	69	0.275	0.401	0.343	0.569	0.646	1.00
A (growth)	OSNet+ctx	69	0.826	0.647	0.739	0.601	0.631	1.00
B (appearance)	Hybrid	82	0.274	0.403	0.342	0.574	0.640	1.00
B (appearance)	OSNet+ctx	82	0.671	0.570	0.663	0.594	0.609	1.00

TABLE II
HYBRID VERSUS THE STRONGEST APPEARANCE BASELINE ON THE SYNTHETIC REVISIT PROTOCOL (SEC. VII-D); PAIRWISE AND ASSIGNMENT F_1 AT THE VALIDATION-CALIBRATED THRESHOLD. ONE SEED (0), TEN SOURCE PHOTOGRAPHS ACROSS FIVE WALLS, NO RESAMPLING — ROWS CARRY NO UNCERTAINTY AND ARE AN ACCURACY CHECKPOINT, NOT A STABILITY CLAIM (SEC. VII-D).

level cost is within an order of magnitude of LoFTR (8.3 s), a crop method with none of the open-set benefit.

D. A synthetic revisit protocol

The maintenance setting these records serve is a *revisited* crack, so we add a second protocol isolating controlled change. Ten source photographs across five of the walls (walls 01, 02, 04, 05, 09; one seed, 0) are each synthesized into a set of queries under an identical warp ($1.08\times$ scale, 6° rotation, 4° tilt): dataset A additionally grows every crack tip (damage evolution), dataset B corrupts appearance while leaving structure untouched (blur $\sigma=1.6$, noise $\sigma=14$, brightness/contrast shift); each synthetic query is ranked against the wall’s real instances only. The *hybrid* folds a structural skeleton comparison into the geometric baseline as a fallback plus per-pair blend, scored by the same code as OSNet+ctx (Table II).

The geometric side does not transfer its main-benchmark advantage here: OSNet+ctx separates the synthetic queries at R@1 0.826/0.671, the hybrid at 0.275/0.274, with open-set rates 0.739/0.663 against 0.343/0.342. But the split is metric-dependent and opposite the ranking gap: on identity agreement the hybrid holds its own — assignment $F_1^a = 0.646/0.640$ beats OSNet’s 0.631/0.609, and pairwise $F_1^{\text{@v}} = 0.569/0.574$ trails 0.601/0.594 by far less than the ranking margin. Both rate 1.00 scored, and per cell none was decided by the pure-homography branch alone ($H\% = 0.0$ — the share of assignment cells whose verdict came straight from the homography, with no blend or skeleton fallback): the geometric pipeline answers *which crack is which* on pairs appearance also resolves, but loses the global ranking of a hundred-plus same-wall candidates at one threshold. These rows are a checkpoint, not a claim: one seed, one wall subset.

VIII. THREATS TO VALIDITY

We state the limits in the body because they bound the central claims.

The effective sample is 19 identities. The test split’s 280 queries come from 19 identities visible in two or more photographs across 11 walls; queries from one wall share material, illumination and one camera pass, so query count is not sample size. We resample walls and report the leave-one-wall-out spread in Rank-1. The same resampling discipline applies to the open-set headline: DIR@FAR=0.1 of 0.768 is a point estimate at one operating point on the same sample, so the abstract compares it to baselines as a one-number structure

— the 0.35 gap — and claims nothing about the precise spread between resampled walls. This benchmark supports a large structural claim — a 0.35 open-set gap, a formulation change that lifts a band nine methods sit inside — but not fine-grained ranking among baselines.

We do not attach a numeric wall-resampled interval to the abstract point estimates because the score matrices needed to recompute it are not part of the released artefacts; the effective sample and wall-level bimodality should therefore be read as a limitation, not as hidden precision.

Scope is controlled, and the control found something. The method reads full photographs; the baselines read crops — dictated by design, since homography needs wall context. Two controls (Sec. V-C) separate identity cue from field of view, and scope does not account for the open-set gap (0.404 best context against 0.768). But the crack-erased control is uncomfortable: inpainting the crack costs only 0.011 of that 0.404, so a context-scope number is partly background-matching, which we expect to weaken on real revisits. Ranking numbers are additionally near-duplicate (the adjacent frame answers all 280 queries, Sec. III); the gap sweep shows the trend is flat, with power falling to 105 answerable queries by gap 3.

Annotation and the geometric baseline share a mechanism. Identities were seeded by propagating one click point between consecutive photographs through an ORB homography (Sec. III-B), and the geometric baseline’s whole mechanism is homography registration — so the ground truth could be biased toward exactly the pairs REG+CHAMFER can answer. This is distinct from the merge-rate disclosure of Sec. III-B: that reports a deliberate design choice, this would be a measurement artefact. On this data the entanglement is real but not separable: of the 118 positive (same-crack) query-gallery pairs on test, 116 (98.3%) involve an identity that was itself formed by a merge, and both-sides-unmerged positives number just 2, so a DIR@FAR restricted to unmerged identities carries no statistical power. Two partial checks mitigate it. First, the mechanism is not exactly shared: annotation warped one ORB homography between *consecutive* frames, while the baseline registers *every* pair with a crack-masked SIFT/MAGSAC/ECC cascade and scores by Chamfer — a strictly harder alignment problem. Second, positives register at the same rate whether or not they involve a merged identity (100% both ways here), so any bias would have to enter through negatives and the decision threshold, not through positive coverage. This check does not test whether a biased homography registers differently from an unbiased one, because both the annotation and baseline alignment are related procedures; nor does it create a clean positive subsample. We report the structure rather than assert freedom from it; a human-audited identification split, with alignment-independent labels, is the clean experiment that would settle it.

Annotations are instrumented, not audited. The human verdicts that would turn the provenance of Sec. III-B into a sampled error rate with a confidence interval have not

been collected; the 45% merge rate is reported, and relative comparison is robust to it since the merge policy hits all methods identically. All captures come from one site over four half-days, two handsets and two surface families; smooth plaster produces the categorical SIFT failure noted in Sec. VII-B; cracks are located by a learned segmenter whose errors propagate into Chamfer agreement unquantified. What would convert the cross-visit claim from inferred to designed — and the synthetic revisits of Sec. VII-D to measurement — is a real second capture session at marked distances and angles; we have not run one, and every multi-visit statement here is a hypothesis, not a result. **Annotation and the geometric baseline share a mechanism.** Identities were seeded by propagating one click point between consecutive photographs through an ORB homography (Sec. III-B), and the geometric baseline’s whole mechanism is homography registration — so the ground truth could be biased toward exactly the pairs REG+CHAMFER can answer. This is distinct from the merge-rate disclosure of Sec. III-B: that reports a deliberate design choice, this would be a measurement artefact. On this data the entanglement is real but not separable: of the 118 positive (same-crack) query-gallery pairs on test, 116 (98.3%) involve an identity that was itself formed by a merge, and both-sides-unmerged positives number just 2, so a DIR@FAR restricted to unmerged identities carries no statistical power. Two partial checks mitigate it. First, the mechanism is not exactly shared: annotation warped one ORB homography between *consecutive* frames, while the baseline registers *every* pair with a crack-masked SIFT/MAGSAC/ECC cascade and scores by Chamfer — a strictly harder alignment problem. Second, positives register at the same rate whether or not they involve a merged identity (100% both ways here), so any bias would have to enter through negatives and the decision threshold, not through positive coverage. We report the structure rather than assert freedom from it; a human-audited identification split is the clean experiment that would settle it.

the effective sample and wall-level bimodality should therefore be read as a limitation, not as hidden precision. We report the structure rather than assert freedom from it; a human-audited identification split, with alignment-independent labels, is the clean experiment that would settle it. **Annotations are**

instrumented, not audited. The human verdicts that would turn the provenance of Sec. III-B into a sampled error rate with a confidence interval have not been collected; the 45% merge rate is reported, and relative comparison is robust to it since the merge policy hits all methods identically. All captures come from one site over four half-days, two handsets and two surface families; cracks are located by a learned segmenter whose errors propagate into Chamfer agreement unquantified. What would convert the cross-visit claim from inferred to designed — and the synthetic revisits of Sec. VII-D to measurement — is a real second capture session at marked distances and angles; we have not run one, and every multi-visit statement here is a hypothesis, not a result.

IX. CONCLUSION

We built the dataset that makes crack re-identification testable — CRACKID, 140 photographs over 17 walls with 212 identities, its annotation provenance (45% merged) stated rather than assumed — and the benchmark’s main finding is negative and structural: on the decision a record requires, nine matchers across four families and two orders of compute buy between 0.000 and 0.054 over chance, and only LoFTR survives a validation-calibrated threshold. Ordering a gallery and deciding an identity are not the same capability. The difficulty is not crack appearance: a geometric reference baseline that registers on wall texture with cracks masked out and Chamfer-matches projected masks reaches DIR@FAR=0.1 of 0.768 against 0.404 for the best context method while answering 48% of pairs — and a control exploiting that alignment coverage without examining a crack scores 0.000. The open-set advantage has a specific, unshared source: an aligned pair supports a *licensed negative*, positive evidence of absence, where a crop embedding can only report low similarity. Two bounds remain addressable — a real second capture session turns the benchmark from viewpoint to cross-visit, and the annotation audit awaits its verdicts. The 48% coverage is also a deployment gap in its own right, not merely a scoring issue: smooth-plaster walls can be categorically unanswerable. Finally, the synthetic revisit protocol reverses the headline ranking and open-set result in favour of OSNet+context, while the geometric pipeline retains only a narrow assignment- F_1 advantage; the present paper therefore establishes a promising within-visit geometric reference, not a solved multi-visit maintenance system.

REFERENCES

- [1] S. Kulkarni, S. Singh, D. Balakrishnan, S. Sharma, S. Devunuri, and S. C. R. Korlapati, “CrackSeg9k: A collection and benchmark for crack segmentation datasets and frameworks,” *arXiv:2208.13054*, 2022.
- [2] C. Benz and V. Rodehorst, “OmniCrack30k: A benchmark for crack segmentation and the reasonable effectiveness of transfer learning,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, 2024.
- [3] K. Zhou, Y. Yang, A. Cavallaro, and T. Xiang, “Omni-scale feature learning for person re-identification,” in *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019, pp. 3702–3712.
- [4] Y. Qian, J. Barthelemy, E. Karuppiyah *et al.*, “Identifying re-identification challenges: Past, current and future trends,” *SN Computer Science*, vol. 5, p. 937, 2024.
- [5] M. J. Hasan, C. K. Nguyen, Y. L. Boo, H. Jahani, and K.-L. Ong, “Vehicle damage detection using artificial intelligence: A systematic literature review,” *WIREs Data Mining and Knowledge Discovery*, vol. 15, p. e70027, 2025.
- [6] S. A. Pérez-Zarate, D. Corzo-García, J. L. Pro-Martín *et al.*, “Automated car damage assessment using computer vision: Insurance company use case,” *Applied Sciences*, vol. 14, no. 20, p. 9560, 2024.
- [7] D. Mallios, L. Xiaofei, N. McLaughlin, J. M. Del Rincon, C. Galbraith, and R. Garland, “Vehicle damage severity estimation for insurance operations using in-the-wild mobile images,” *IEEE Access*, vol. 11, pp. 78 644–78 655, 2023.
- [8] L. Parente, E. Falvo, C. Castagnetti, F. Grassi, F. Mancini, P. Rossi, and A. Capra, “Image-based monitoring of cracks: Effectiveness analysis of an open-source machine learning-assisted procedure,” *Journal of Imaging*, vol. 8, no. 2, p. 22, 2022.
- [9] S. Harb, P. Achanccaray, M. Maboudi, and M. Gerke, “Multi-temporal crack segmentation in concrete structures using deep learning approaches,” *ISPRS Annals of the Photogrammetry, Remote Sensing and Spatial Information Sciences*, vol. X-G-2025, pp. 341–348, 2025.
- [10] X. Sun and P. Chang, “Robust perspective correction for real-world crack evolution tracking in image-based structural health monitoring,” *arXiv:2506.22437*, 2025.
- [11] NEC Corporation, “Photographing guide device,” US Patent 12,223,639 B2 (priority 2019), 2025.
- [12] S. Liao, Z. Mo, Y. Hu, and S. Z. Li, “Open-set person re-identification,” in *arXiv:1408.0872*, 2014.
- [13] D. G. Lowe, “Distinctive image features from scale-invariant keypoints,” *International Journal of Computer Vision*, vol. 60, no. 2, pp. 91–110, 2004.
- [14] E. Rublee, V. Rabaud, K. Konolige, and G. Bradski, “ORB: An efficient alternative to SIFT or SURF,” in *IEEE International Conference on Computer Vision (ICCV)*, 2011, pp. 2564–2571.
- [15] P.-E. Sarlin, D. DeTone, T. Malisiewicz, and A. Rabinovich, “SuperGlue: Learning feature matching with graph neural networks,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 4938–4947.
- [16] J. Sun, Z. Shen, Y. Wang, H. Bao, and X. Zhou, “LoFTR: Detector-free local feature matching with transformers,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 8922–8931.
- [17] A. Dosovitskiy, L. Beyer, A. Kolesnikov *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *International Conference on Learning Representations (ICLR)*, 2021.
- [18] H. Touvron, M. Cord, M. Douze, F. Massa, A. Sablayrolles, and H. Jégou, “Training data-efficient image transformers and distillation through attention,” in *International Conference on Machine Learning (ICML)*, 2021, pp. 10 347–10 357.
- [19] A. Radford, J. W. Kim, C. Hallacy *et al.*, “Learning transferable visual models from natural language supervision,” in *International Conference on Machine Learning (ICML)*, 2021, pp. 8748–8763.
- [20] M. Oquab, T. Darcet, T. Moutakanni *et al.*, “DINOv2: Learning robust visual features without supervision,” *Transactions on Machine Learning Research*, 2024.
- [21] L. Maiano, A. Montuschi, M. Caserio *et al.*, “A deep-learning-based antifraud system for car-insurance claims,” *Expert Systems with Applications*, vol. 231, p. 120644, 2023.
- [22] K. Sakurada and T. Okatani, “Change detection from a street image pair using CNN features and superpixel segmentation,” in *British Machine Vision Conference (BMVC)*, 2015.
- [23] P. F. Alcantarilla, S. Stent, G. Ros, R. Arroyo, and R. Gherardi, “Street-view change detection with deconvolutional networks,” *Autonomous Robots*, vol. 42, no. 7, pp. 1301–1322, 2018.
- [24] R. Hartley and A. Zisserman, *Multiple View Geometry in Computer Vision*, 2nd ed. Cambridge University Press, 2004.
- [25] M. A. Fischler and R. C. Bolles, “Random sample consensus: A paradigm for model fitting with applications to image analysis and automated cartography,” in *Readings in Computer Vision*. Morgan Kaufmann, 1987, pp. 726–740.
- [26] D. Barath, J. Matas, and J. Noskova, “MAGSAC: Marginalizing sample consensus,” in *IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 10 197–10 205.
- [27] G. D. Evangelidis and E. Z. Psarakis, “Parametric image alignment using enhanced correlation coefficient maximization,” *IEEE Trans. Pattern Analysis and Machine Intelligence*, vol. 30, no. 10, pp. 1858–1865, 2008.